Optimizing Decision-Making in Multi-Agent RL with CPT

Sheyan Lalmohammed, Khush Gupta, Alok Shah

Introduction

- Investigating Multi-Agent Reinforcement Learning (MARL) under Cumulative Prospect Theory (CPT)
- **Motivation:** Aligning autonomous agents with human decision-making biases
- **Key Questions:**
  - Do CPT-trained agents adhere to their utility and probability distortion functions?
  - How do CPT-guided agents optimize strategies in multi-agent games compared to those using traditional utility functions?
  - To what extent do agents adapt their strategies based on the utility functions of counterparties? What emergent dynamics arise in mixed populations of agents?

Background on CPT - Prospect Theory

Developed by Daniel Kahneman and Amos Tversky in 1979, Prospect Theory explains decision-making under risk and uncertainty:

- **Loss Aversion:** People tend to avoid losses more than acquiring equivalent gains.
- **Relative Evaluation:** Decisions are made based on relative differences rather than absolute outcomes.
- **Reference-Dependent:** Evaluations are based on outcomes relative to a reference point.

Background on CPT - Development of CPT

Cumulative Prospect Theory (CPT) extends the original framework to better handle multiple outcome probabilities:

- **Probability Weighting Function:** Captures the tendency to overweight small probabilities and underweight large probabilities.
- **Value Function:** Typically concave for gains and convex for losses, reflecting human risk preferences.

Implementation Strategy

Technical Approach & PyTorch Implementation

- **Policy Gradient Optimization with CPT:**
  - Integrates CPT-adjusted rewards and probability distortions.
  - Employs model-free learning using policy gradients.
- **Implementation Workflow:**
  i. Design a neural network for policy representation.
  ii. Transform rewards using CPT functions.
  iii. Compute policy gradients via automatic differentiation.
  iv. Optimize policies using gradient ascent.

CPT-Adjusted Rewards

(1)

$$C(X) \;=\; \int_{-\infty}^{0} w^{+}\big(P(u(X) > z)\big)\,dz \;-\; \int_{0}^{\infty} w^{-}\big(P(u(X) > z)\big)\,dz.$$

(2)

$$\max_{\pi \in \Pi_{M,N}} \; C\Big(\sum_{t=0}^{H-1} r_t\Big).$$

- C: The CPT value capturing the decision-maker's distortion in perceiving gains and losses.
- $\pi \in \Pi_{M,N}$: A policy π chosen from the set $\Pi_{M,N}$ (e.g., all feasible memory-based policies).

Algorithms + Policy Gradient

Below is the Policy Gradient that we use to optimize the policy and solve our optimization problem.

$$\nabla_\rho J \;=\; \mathbb{E}\Big[\,\xi\big(\textstyle\sum_\tau r_\tau\big)\,\nabla_\rho\,\mu_\tau\big(a_\tau \mid Q_\tau(s_\tau, a_\tau; n)\big)\,\nabla_{a_\tau} Q_\tau\big(s_\tau, a_\tau; n\big)\Big].$$

$$\xi(V) \;=\; \int_0^{\max(V,0)} w^{+}\Big(\,P\big(u\big(\textstyle\sum_\tau r_\tau\big) > z\big)\Big)\,dz \;-\; \int_0^{\max(-V,0)} w^{-}\Big(\,P\big(u\big(\textstyle\sum_\tau r_\tau\big) > z\big)\Big)\,dz.$$

The algorithm we use is the Multi Agent Deep Deterministic Policy Gradient (MADDPG)

Competitive Environment - Overview

PettingZoo's **Simple Tag** Environment is a basic Multi-Agent Particle Environment (MPE) designed for competition between agents

- **Objective**: Predators work to “tag” or catch the prey, while the prey’s goal is to evade capture.
- **Rewards**: Rewards are structured so that predators gain rewards when they successfully tag the prey, and the prey receives a penalty when caught.

Competitive Environment - Rewards

Baseline

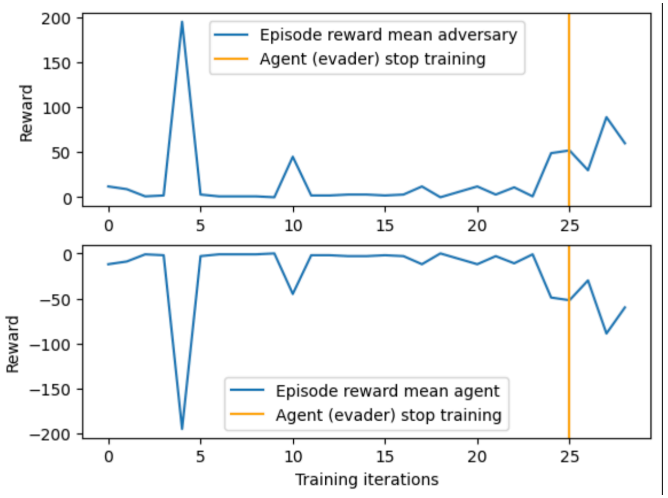

Moderate CPT (Risk Seeking)

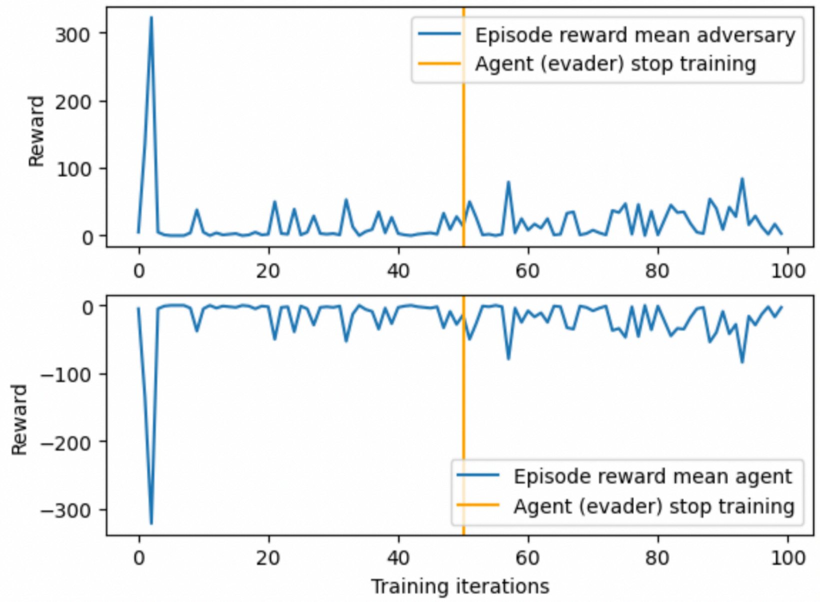

Extreme CPT (Risk Averse)

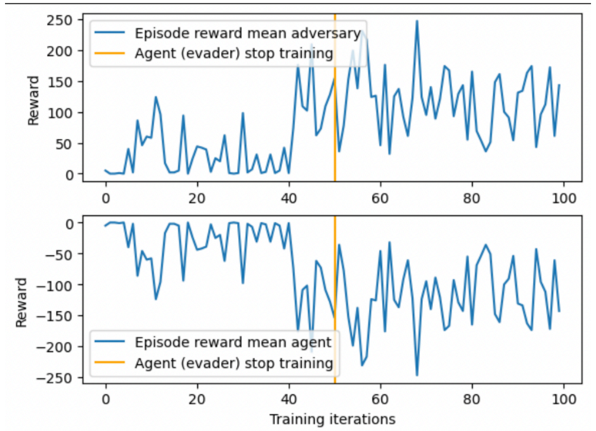

Cooperative Environment - Overview

PettingZoo's **Simple Spread** Environment is a basic Multi-Agent Particle Environment (MPE) designed for semi-collaboration between agents

- **Objective**: The agents work cooperatively to cover all the landmarks. Their goal is to position themselves so that each landmark is “covered” by at least one agent, maximizing overall performance.
- **Rewards**: Rewards encourage efficient coverage of landmarks while also penalizing agents for collisions with one another, which promotes coordinated movement and spacing.

Cooperative Environment - Rewards

Baseline

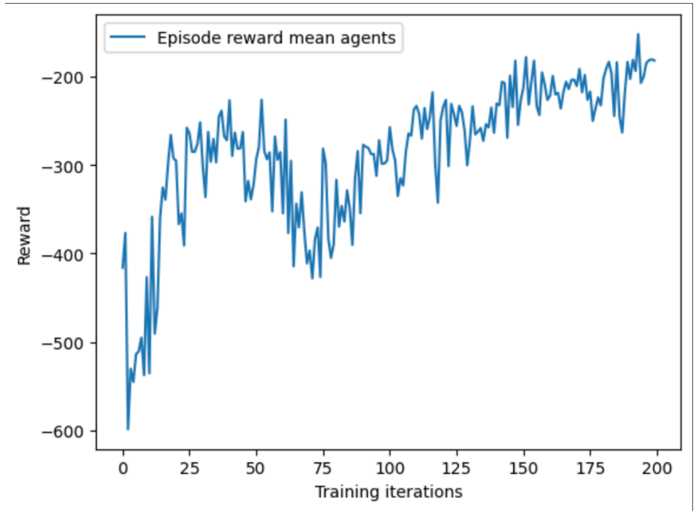

Moderate CPT

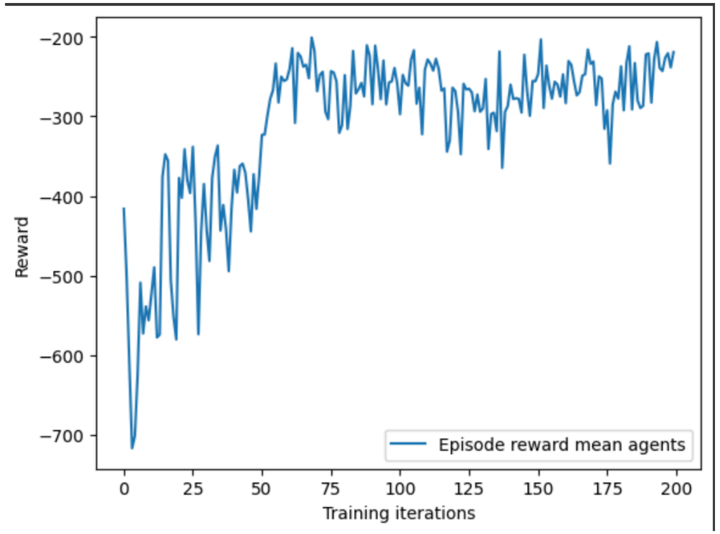

Extreme CPT

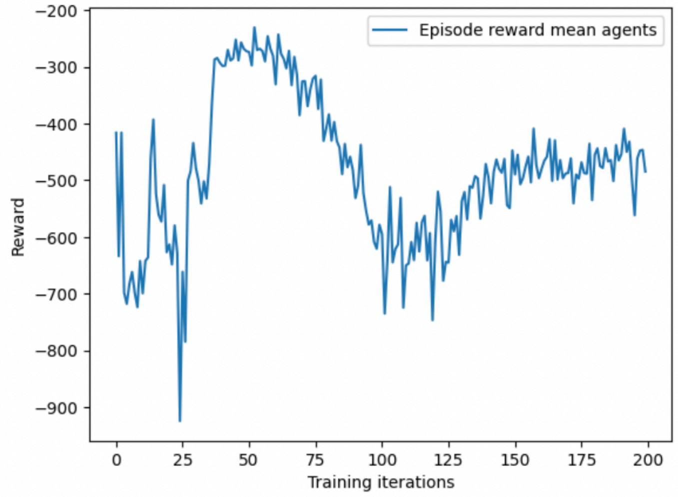

Cooperative Environment - Visualization of MPE

Baseline CPT	Extreme CPT
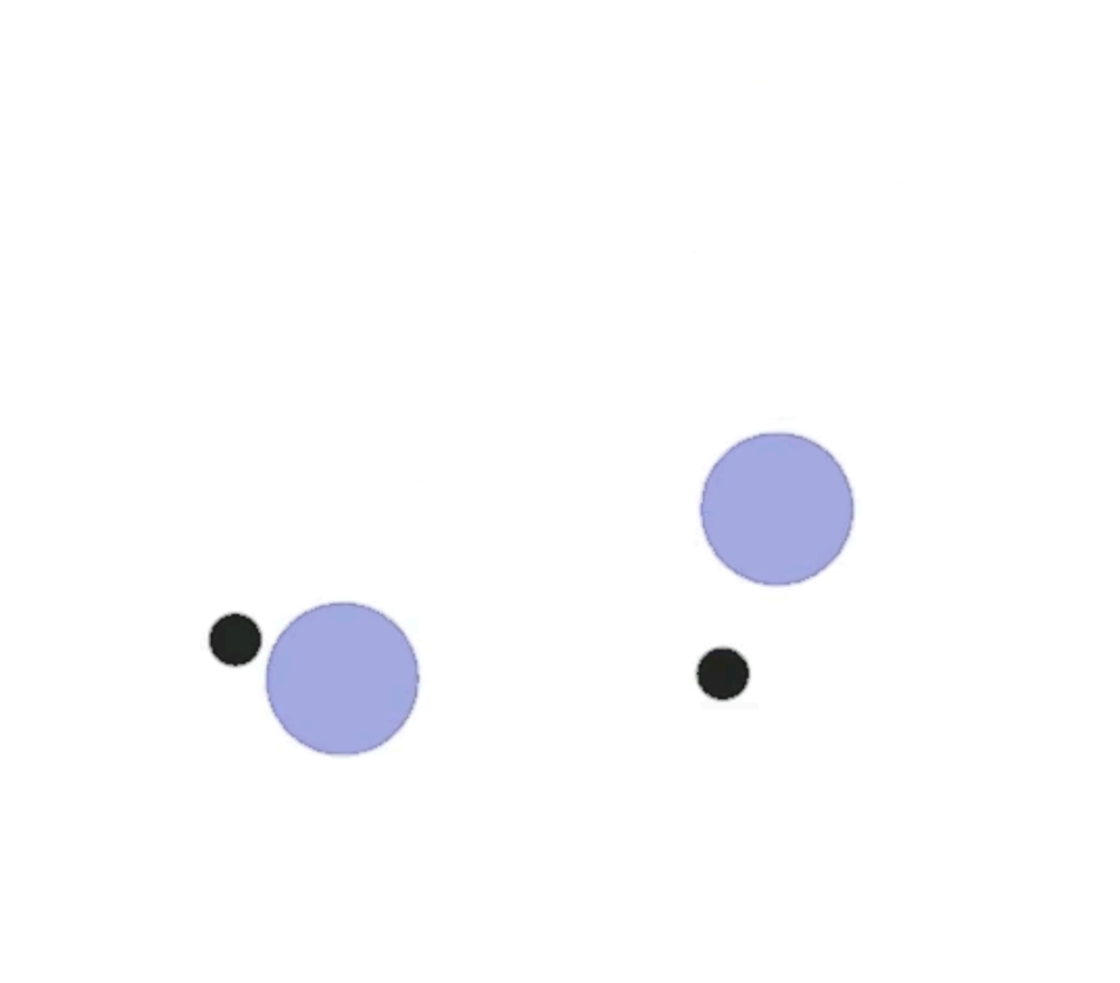	

Cooperative Environment - Dynamic

Risk Seeking

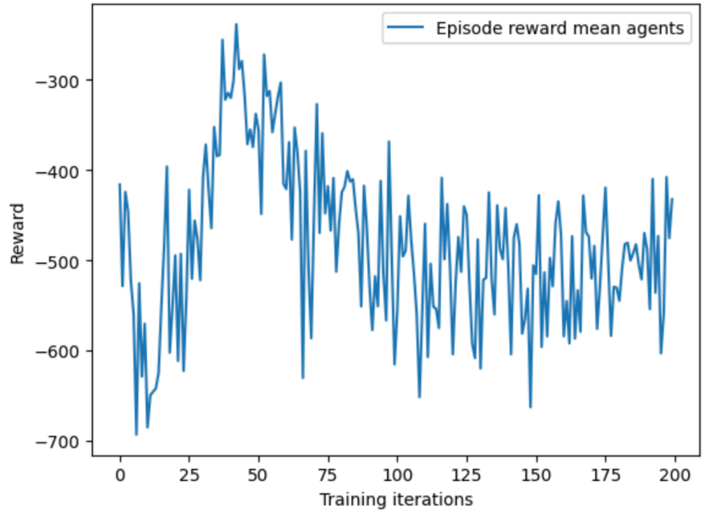

Moderate CPT

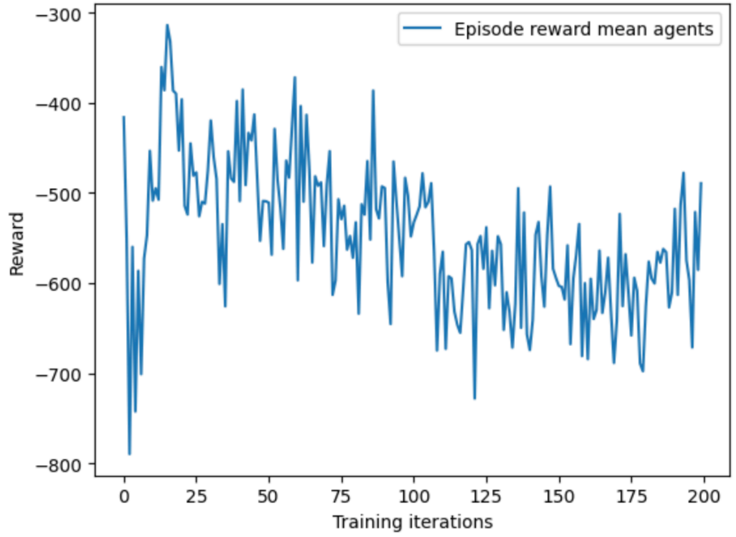

Low CPT

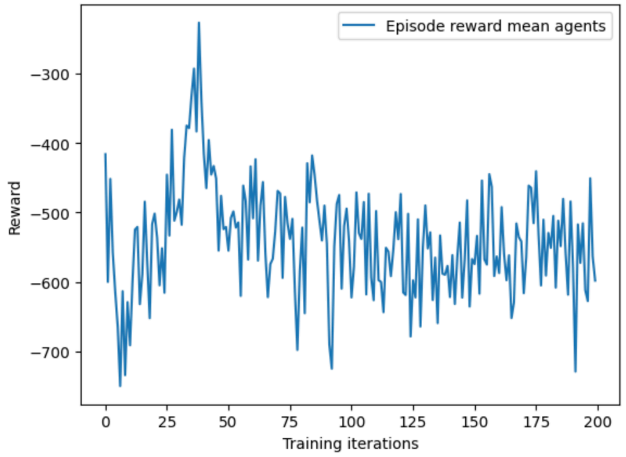

Cooperative Environment - Testing

Seeing - Risk Averse

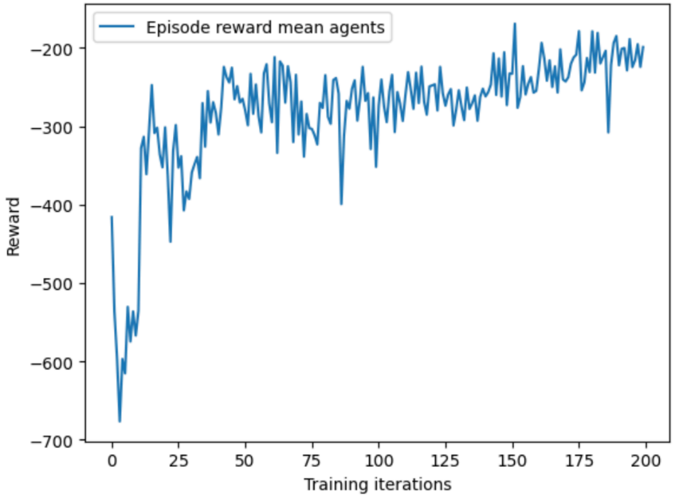

Seeing - Risk Averse and Risk Seeking

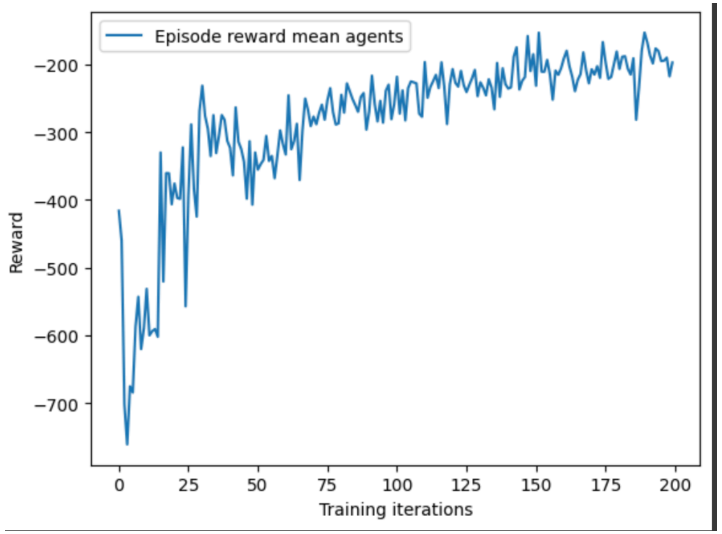

Not Seeing

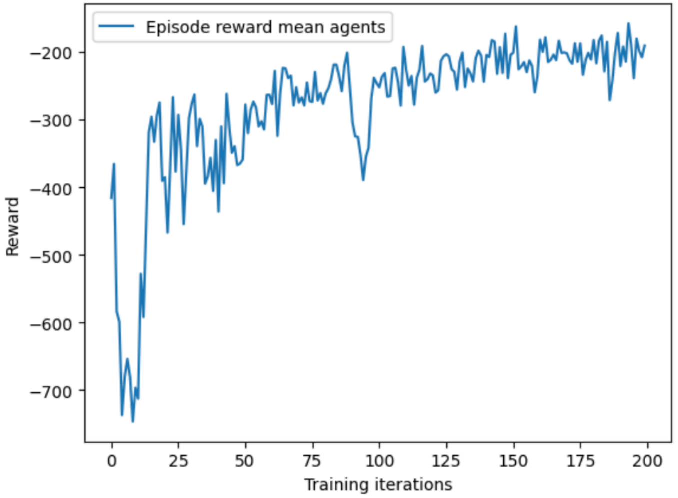

Next Steps & Challenges

Planned Improvements

- **Optimizing CPT Integration**
  - Attempt to try new probability weighting and value distortions
  - See the effect of new estimation methods for the value functions and integral
- **Implementing Discrete Competitive Environments**
  - Try the effect of the CPT-driven policy on an environment like Poker
  - Attempt to induce more interpretable CPT effects driven by Behavioral Economics Studies

Conclusion

Thank You! Questions?